

Parallel/GPU Computing and HPC Research: Past, Present and Future



Research Conclave 2025 Talk @ TCE

Rajesh Pandian M

www.mrprajesh.co.in

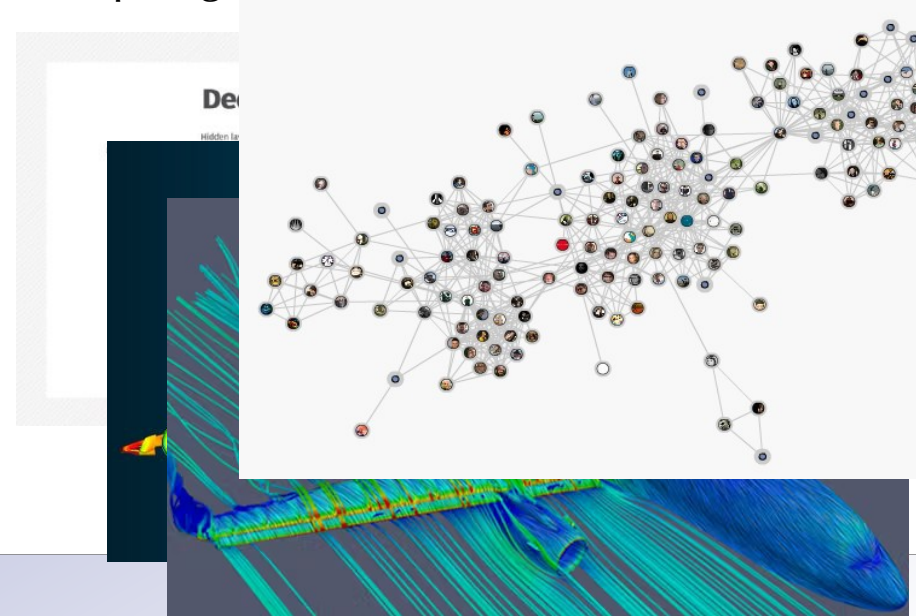

Outline

- Motivation
- Beginning of GPGPU era
 - It all started here in India/IIT
- Present
 - Graph algorithms
 - AI/ML Training
- Future
- Summary

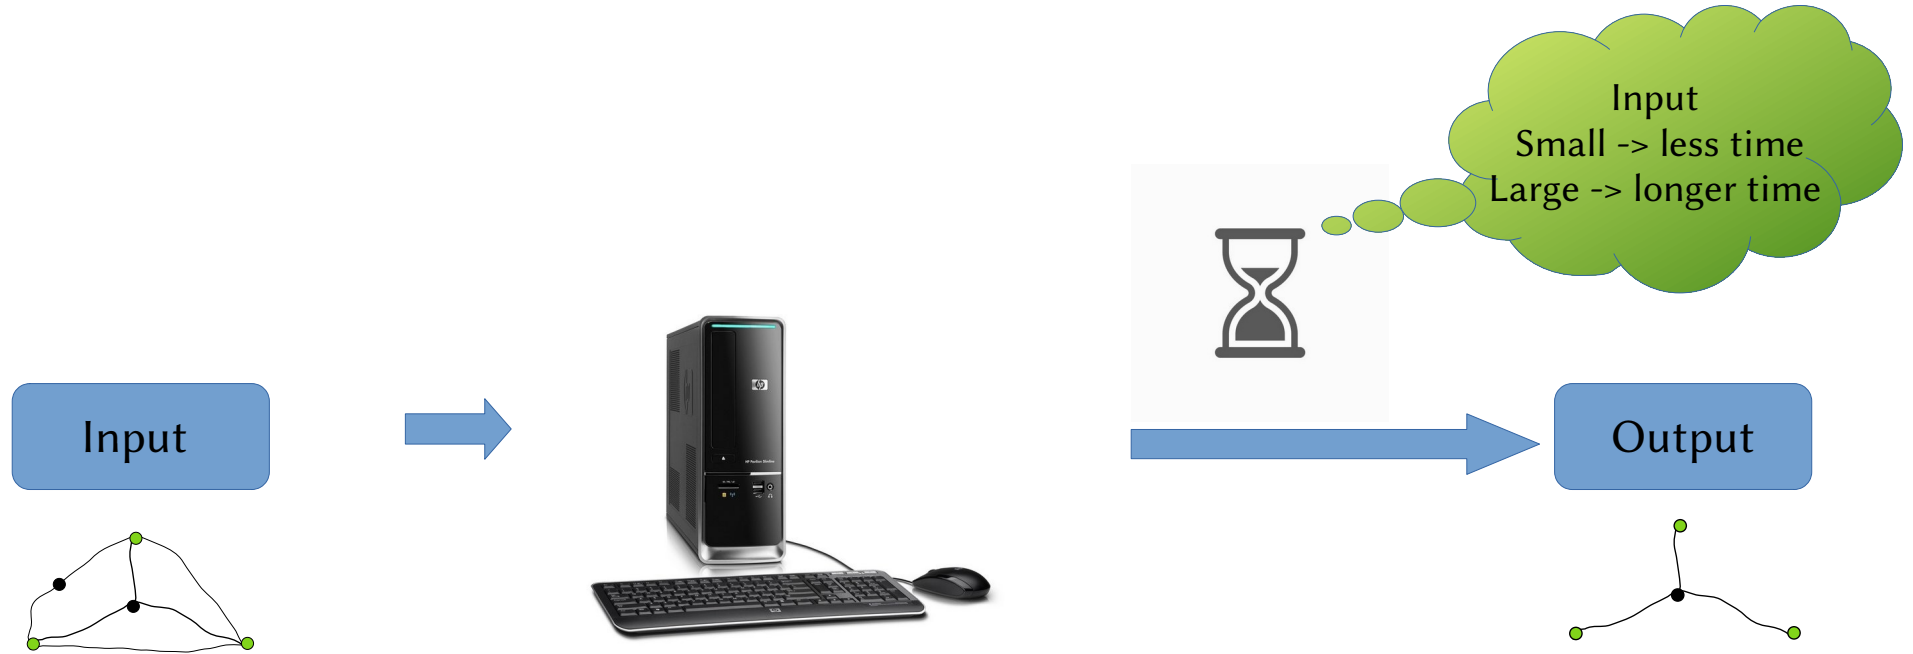


Motivation

- Graphics
 - 3D Games
 - Video rendering
 - 3D Graphics
 - Animations
- Non-Graphics
 - Machine learning
 - Scientific simulations
 - HPC
 - Graph algorithms



Sequential Computing

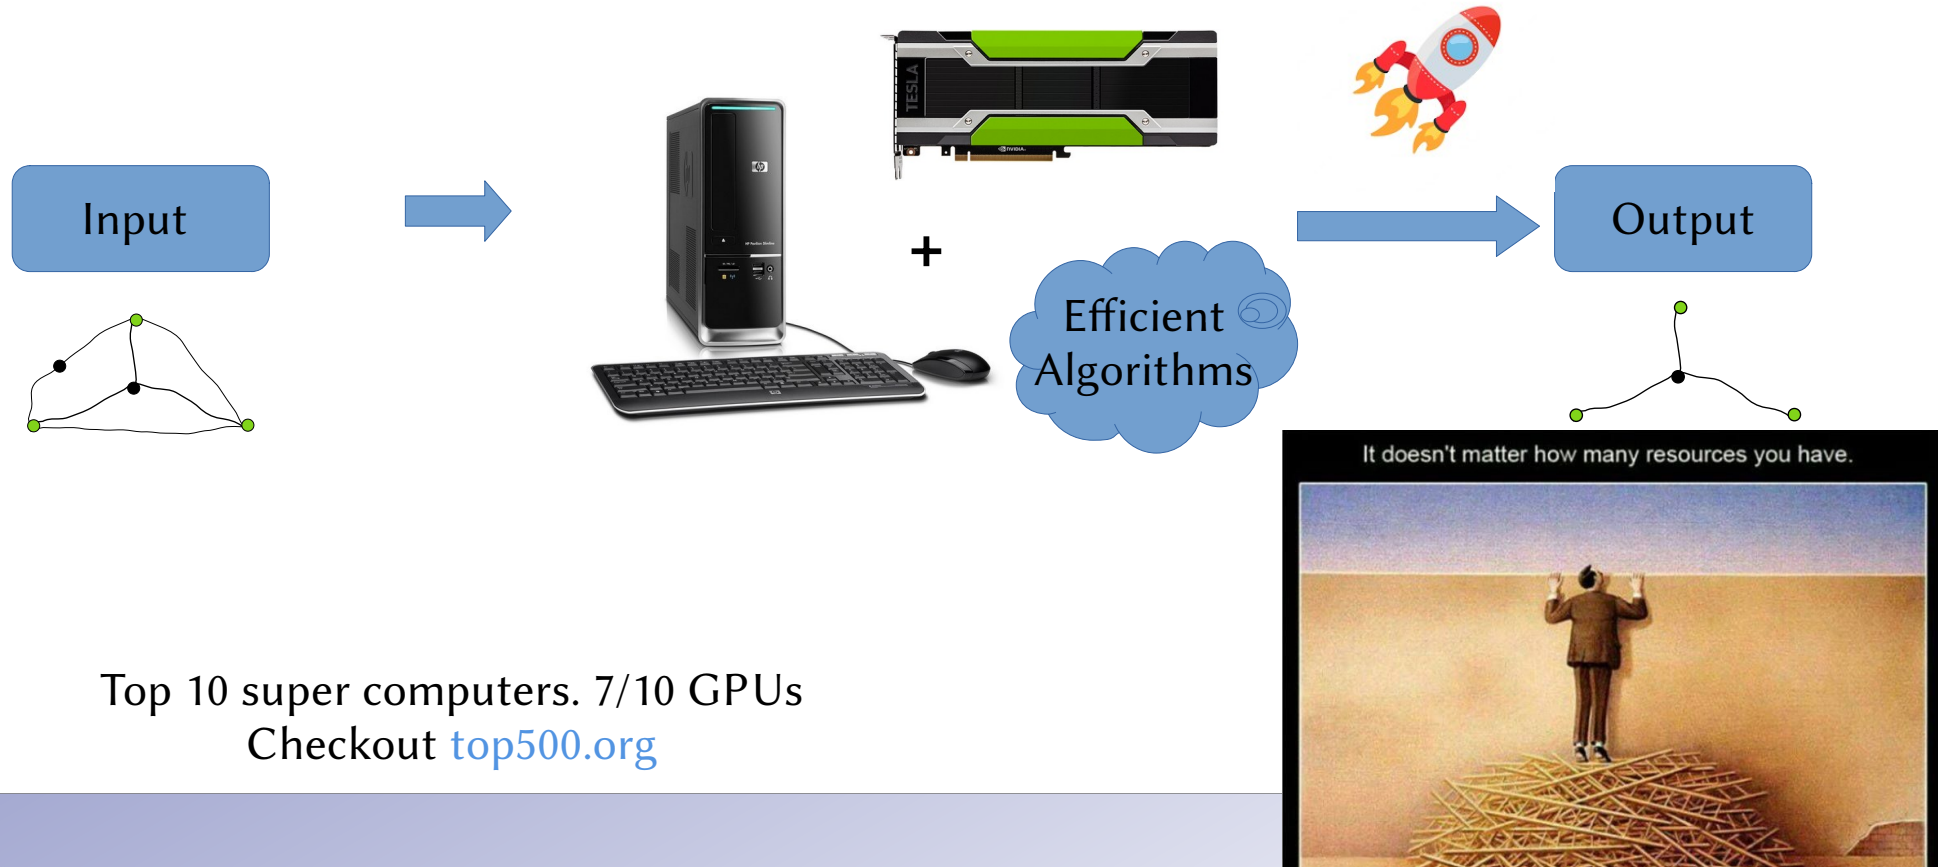


Graph

- Minimum Spanning Tree
- Single source shortest path



Parallel Computation



2016. Joel Hruska 10 years ago, Geforce 8800 graphics card changed PC gaming, computing forever



Parallel Platforms

Software/Frameworks

OpenMP

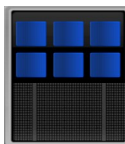
ROCm

MPI


OpenACC

OpenCL

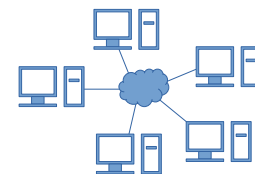
Underlying hardware for parallelization



CPU



GPU



Distributed



The beginning of GPGPU

Accelerating large graph algorithms on the GPU using CUDA

Pawan Harish and P. J. Narayanan

Center for Visual Information Technology
International Institute of Information Technology Hyderabad, INDIA
{harishpk@research., pjn@}iiit.ac.in

HiPC'2007

Abstract. Large graphs involving millions of vertices are common in many practical applications and are challenging to process. Practical-time implementations using high-end computers are reported but are accessible only to a few. Graphics Processing Units (GPUs) of today have high computation power and low price. They have a restrictive programming model and are tricky to use. The G80 line of Nvidia GPUs can be treated as a SIMD processor array using the CUDA programming model. We present a few fundamental algorithms – including breadth first search, single source shortest path, and all-pairs shortest path – using CUDA on large graphs. We can compute the single source shortest path on a 10 million vertex graph in 1.5 seconds using the Nvidia 8800GTX GPU costing \$600. In some cases optimal sequential algorithm is not the fastest on the GPU architecture. GPUs have great potential as high-performance co-processors.



Why GPU for graph processing?

- GPU for General Purpose = GPGPU

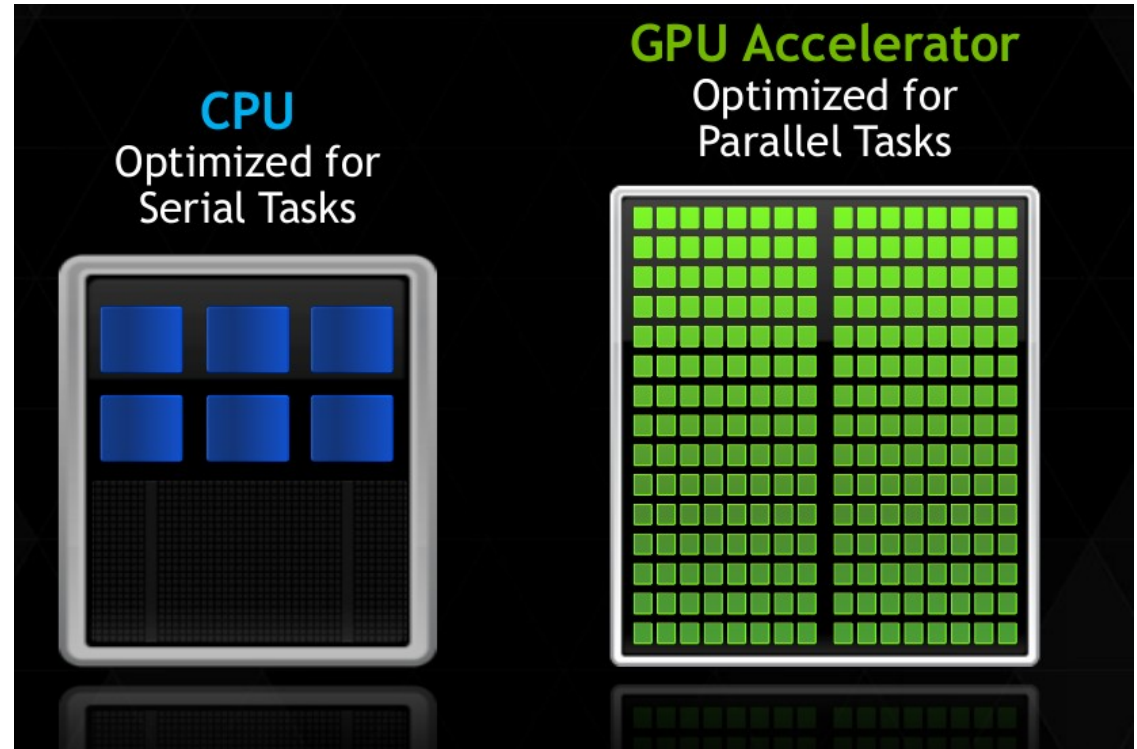
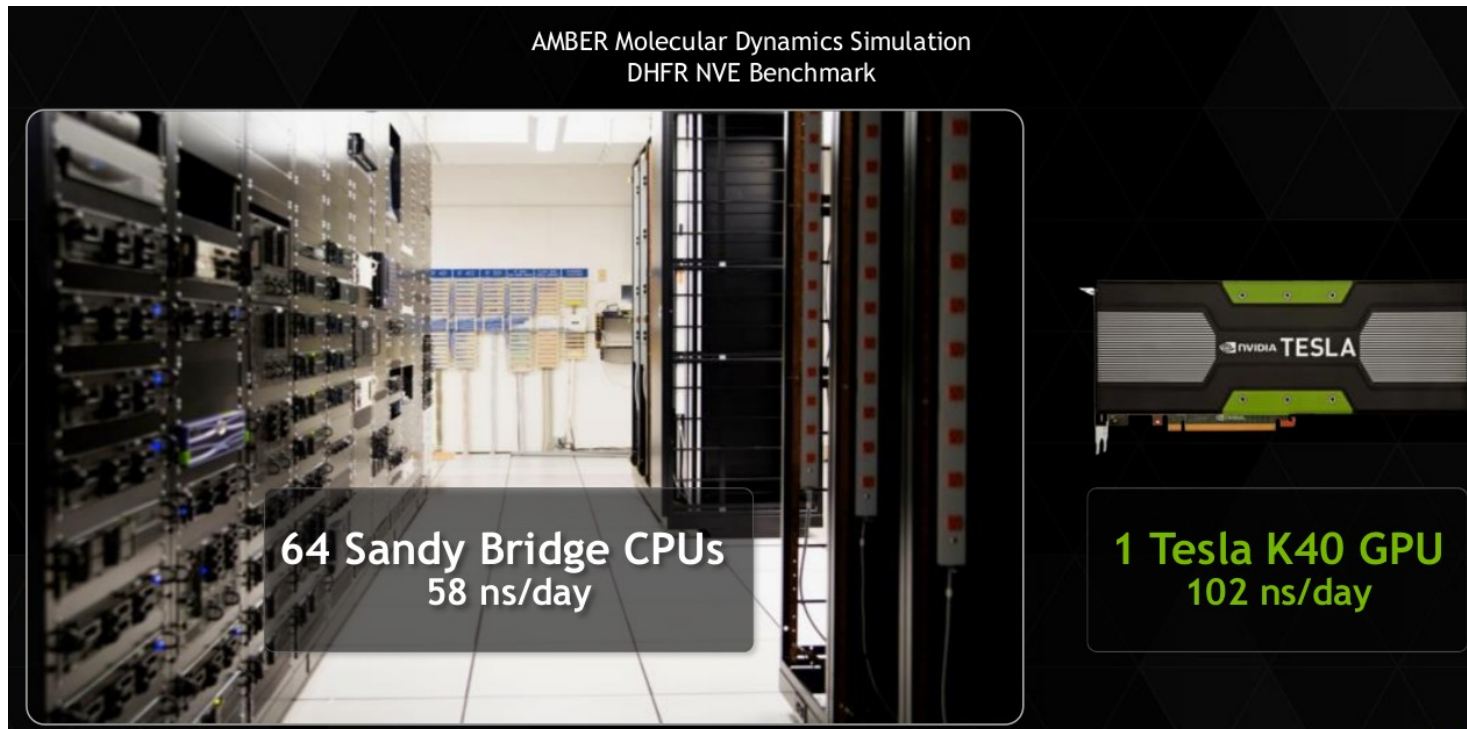


Image Source: GPU Technology Conference (GTC) 2014



Why GPU for graph processing?

- GPU for General Purpose = GPGPU
- **Computation power**



Source: GPU Technology Conference (GTC) 2014
Harvey et al. *Journal of Chemical Theory and Computation*, 2009

Why GPU for graph processing?

- GPU for General Purpose = GPGPU
- Computation power
- Lesser cost
- Energy efficient compared with CPUs

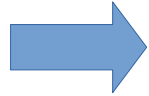


Image Source: GPU Technology Conference (GTC) 2014

Graphs are everywhere

- Networks: Road, Social, Biological, Wireless, and more.
- Such Graphs are extremely large.
- Faster computation is expected.


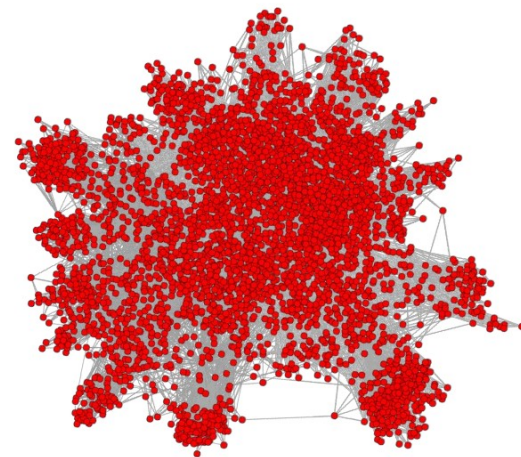
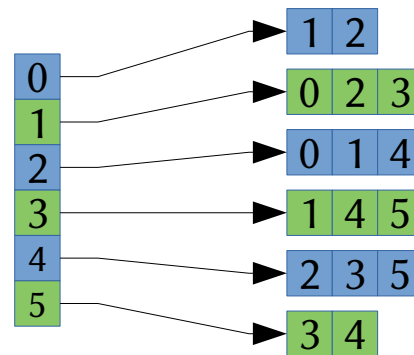
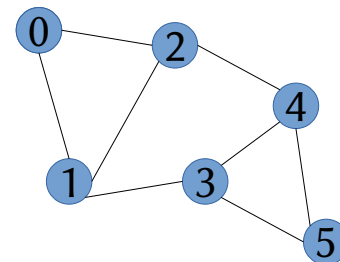


Image Source: <https://spcl.inf.ethz.ch/Publications/.pdf/pushpull-slides.pdf>

Graph Representation

- Adjacency Matrix
- Adjacency list

0	0	1	1	0	0	0
1	1	0	1	1	0	0
2	1	1	0	0	1	0
3	0	1	0	0	1	1
4	0	0	1	1	0	1
5	0	0	0	1	1	0
	0	1	2	3	4	5




Graph Representation – EdgeList

- **EdgeList / Coordinate Format (COO)**
- Compressed Spare Row (CSR)

Reduce repetitions?

0 0 1 1 1 2 2 2 3 3 3 4 4 4 5 5

1 2 0 2 3 0 1 4 1 4 5 2 3 5 3 4


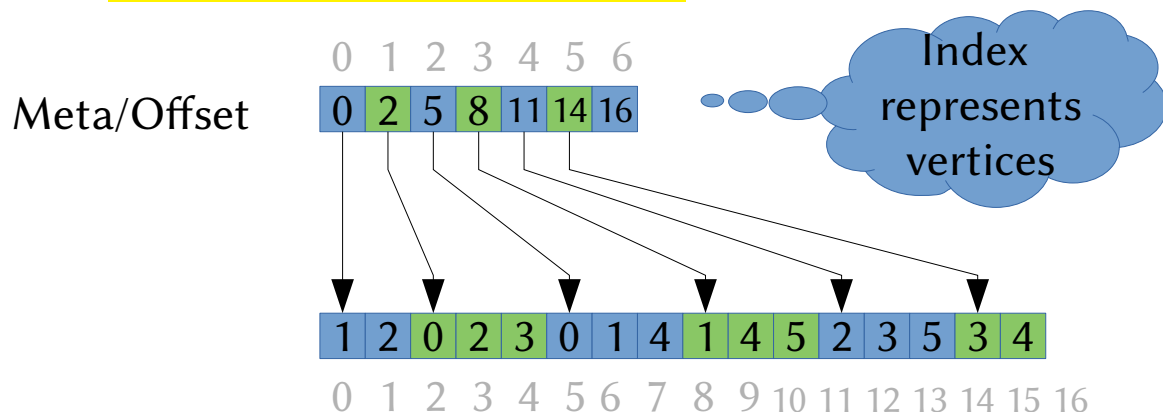


Two 2m-sized array



Graph Representation - CSR

- EdgeList / Coordinate Format (COO)
- **Compressed Spare Row (CSR)**



Two arrays:

Offset array : $n+1$

CSRList : $2m$

Well suited
on GPUs

<https://networkrepository.com/mtx-matrix-market-format.html>



Operator-formulation approach

```
Initialize(G,...);  
do {  
    Operator-Apply(G,...);  
    // on set of nodes, edges  
    // or property of G  
} while (condition);
```

Parallel

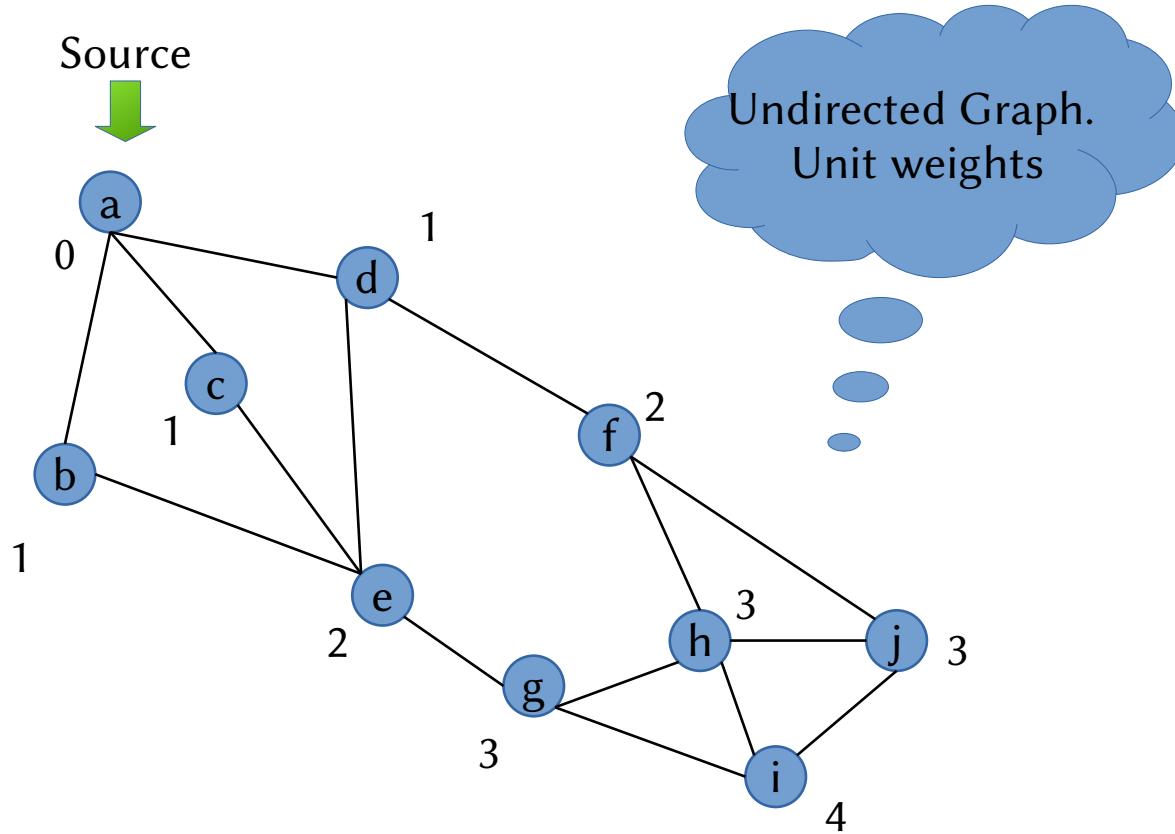
Operator = GPU kernel

- BFS // Computing level information from source
- SSSP // Propagating distance from source




Pingali et. al. *The Tao of parallelism in algorithms*. PLDI, 2011

BFS/SSSP



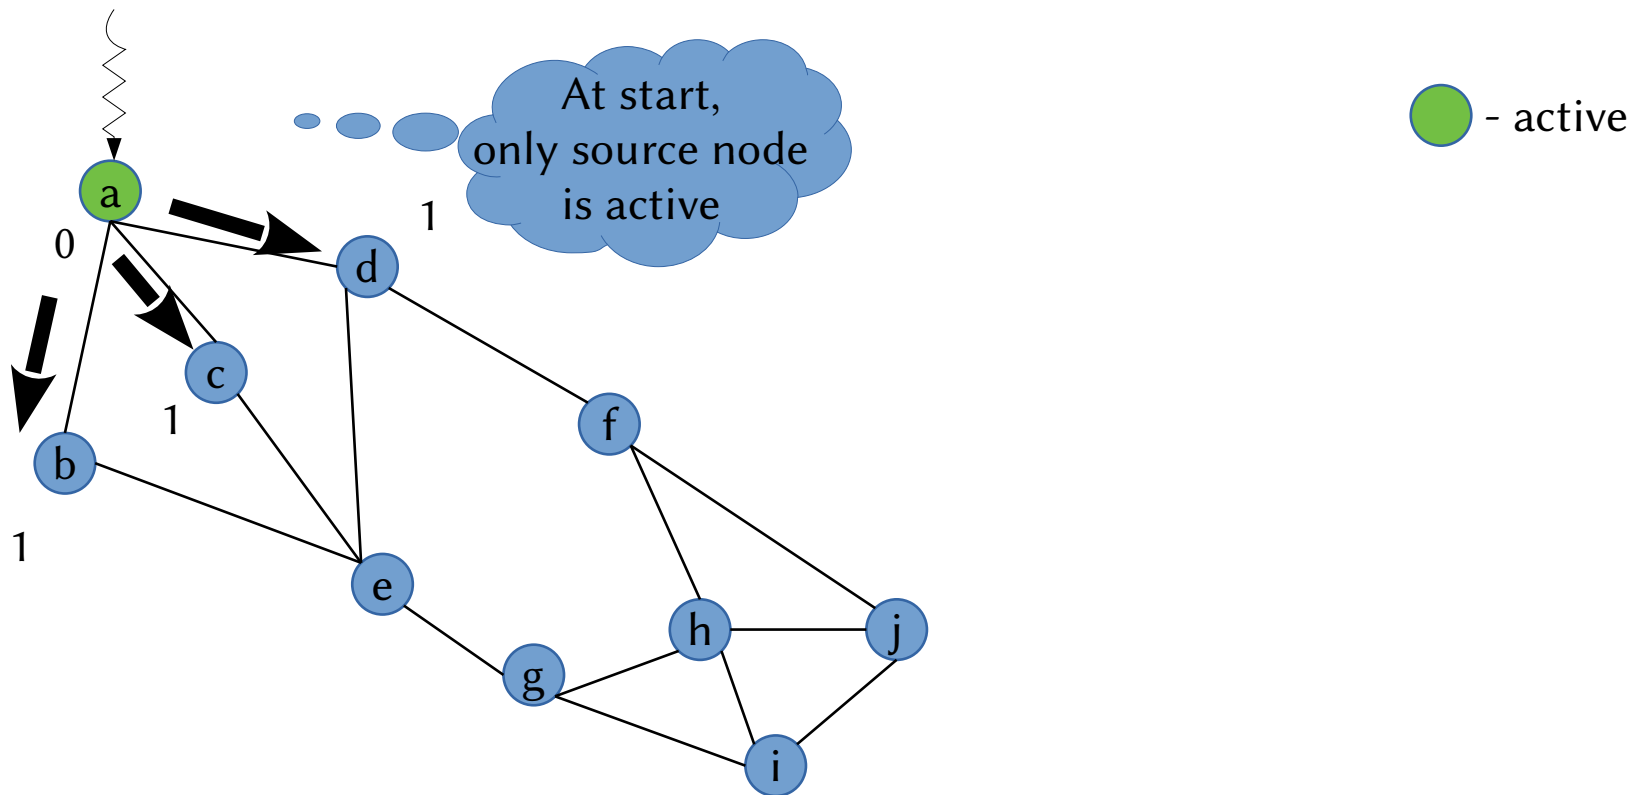
SSSP: Single Source Shortest Path



Rupesh Nasre; Martin Burtscher; Keshav Pingali.

Data-Driven Versus Topology-driven Irregular Computations on GPUs, IPDPS 2013.

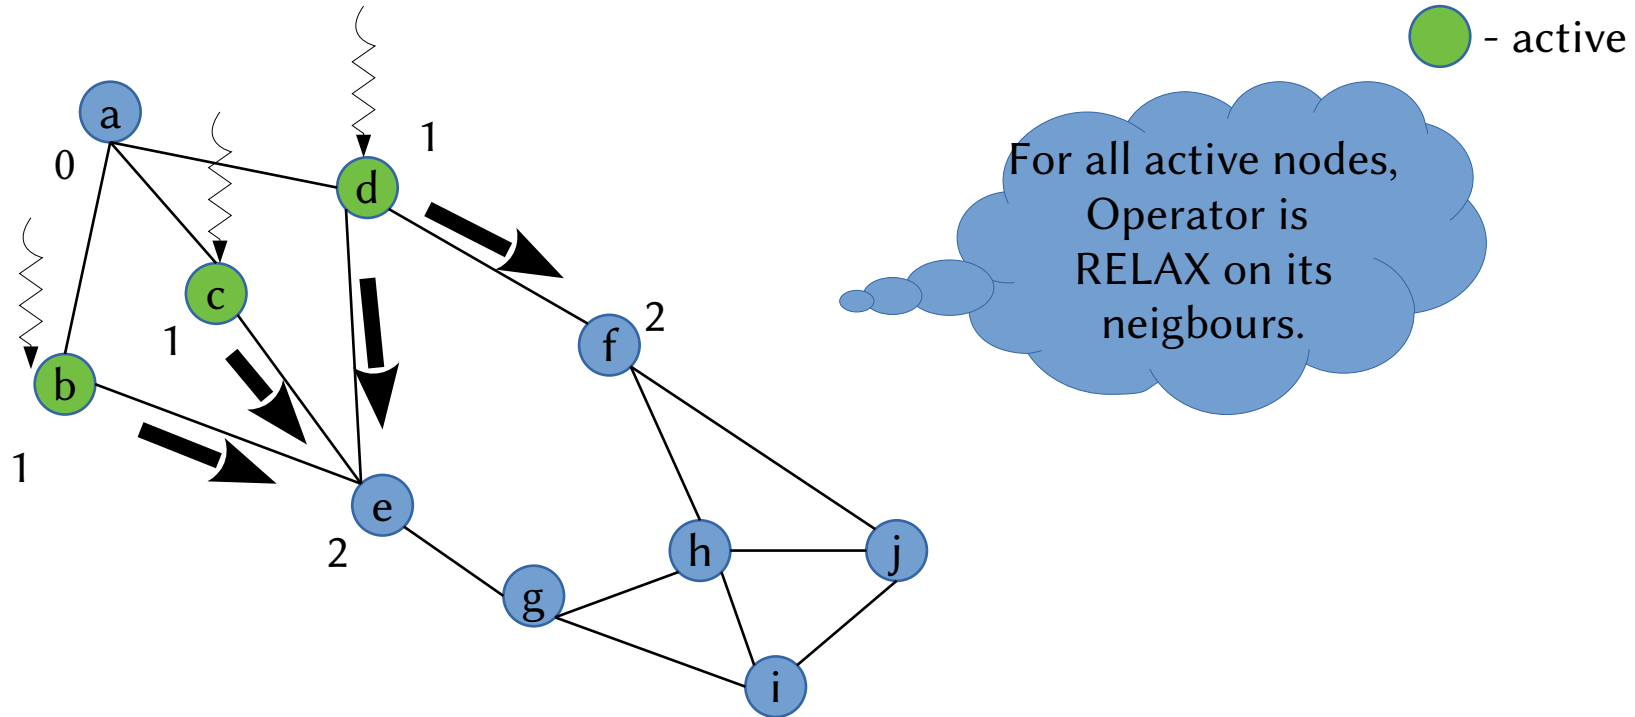
SSSP: Data-driven



Rupesh Nasre; Martin Burtscher; Keshav Pingali.

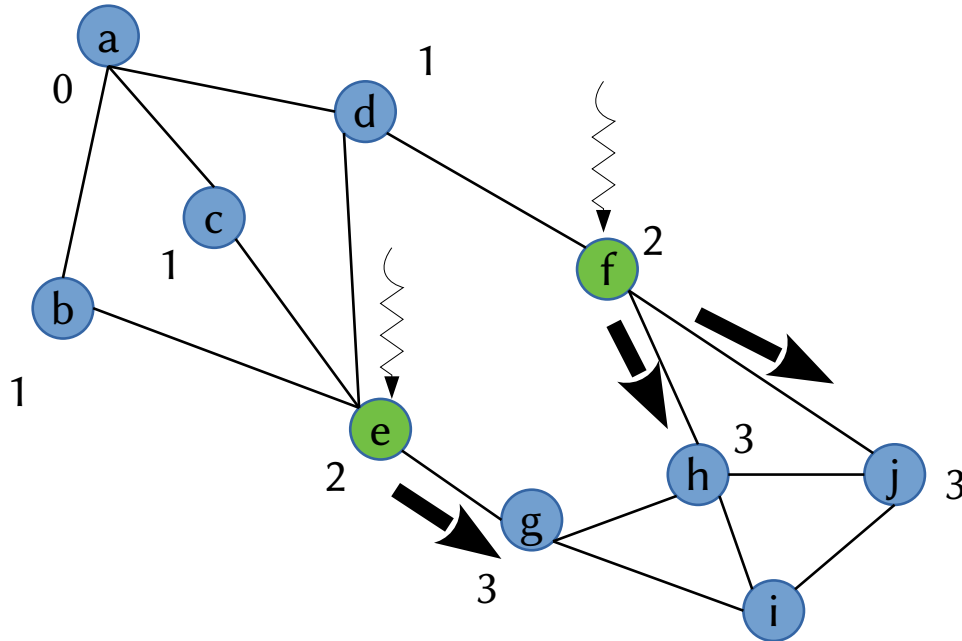
Data-Driven Versus Topology-driven Irregular Computations on GPUs, IPDPS 2013.

SSSP: Data-driven



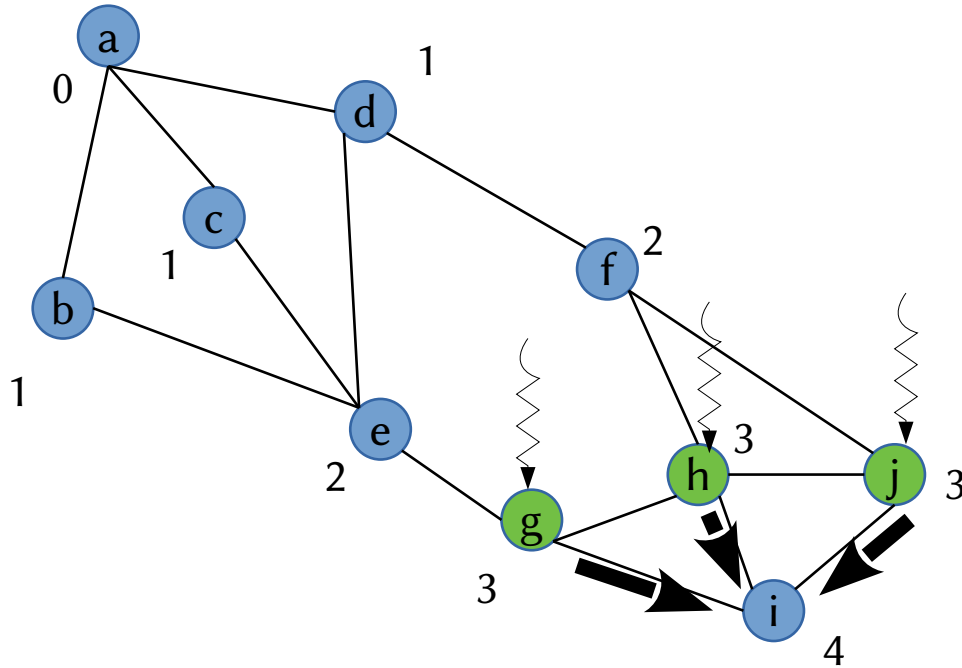
SSSP: Data-driven

● - active



SSSP: Data-driven

● - active




- Stop condition = No active node
- This Algorithm is Poly-time





DEPARTMENT OF COMPUTER SCIENCE & ENGINEERING
INDIAN INSTITUTE OF TECHNOLOGY MADRAS
CHENNAI – 600036

NP-hard Problems meet Parallelization



A Thesis

Submitted by

RAJESH PANDIAN M

Steiner Tree Problem

Generalization of Min Spanning Tree

Vehicle Routing

Generalization of Travelling Salesman

Best sequential
Algorithm



Run on
GPU/Parallel HW



Our Philosophy



... take a **fresh look** at some of the classic graph algorithms and devise **faster** and more parallel GPU and CPU implementations.

+

NP-hard

=

Our Philosophy

- Fallin et al.

A High-Performance MST Implementation for GPUs

Alex Fallin
Dept. of Computer Science
Texas State University
San Marcos, Texas, USA
waf13@txstate.edu

Andres Gonzalez
Dept. of Computer Science
Texas State University
San Marcos, Texas, USA
ag1548@txstate.edu

Jarim Seo
Dept. of Computer Science
Texas State University
San Marcos, Texas, USA
j_s1195@txstate.edu


Martin Burtscher
Dept. of Computer Science
Texas State University
San Marcos, Texas, USA
burtscher@txstate.edu

SC'23

ABSTRACT

Finding a minimum spanning tree (MST) is a fundamental graph algorithm with applications in many fields. This paper presents ECL-MST, a fast MST implementation designed specifically for GPUs. ECL-MST is based on a parallelization approach that unifies Kruskal's and Borůvka's algorithm and incorporates new and existing optimizations from the literature, including implicit path compression and edge-centric operation. On two test systems, it outperforms leading GPU and CPU codes from the literature on all of our 17 input graphs from various domains. On a Titan V GPU,


lines. In this example, the cheapest distribution grid that allows everyone to deliver or receive electricity is the MST shown.



Optimizing for peak performance

OPTIMIZATIONS

- Input Characteristics
- Properties of substeps



- Vertex-/Edge- centric
- Data-/Topology-driven
- Push-/Pull-based

...

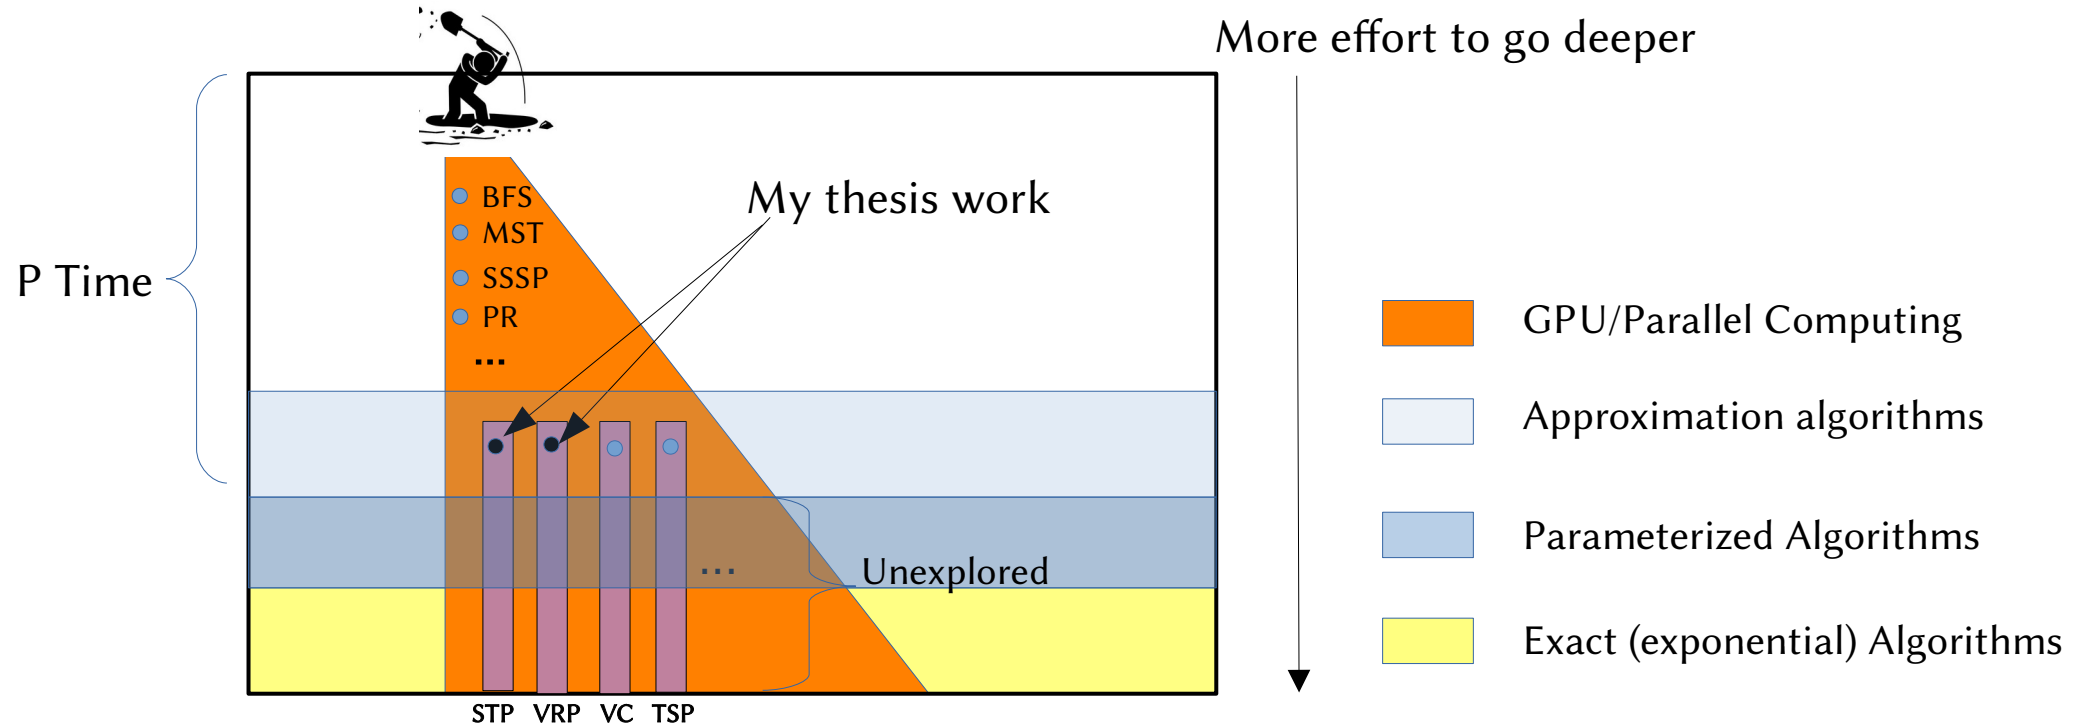
Even more
Peak performance

- Shared memory
- Warp-Intrinsics

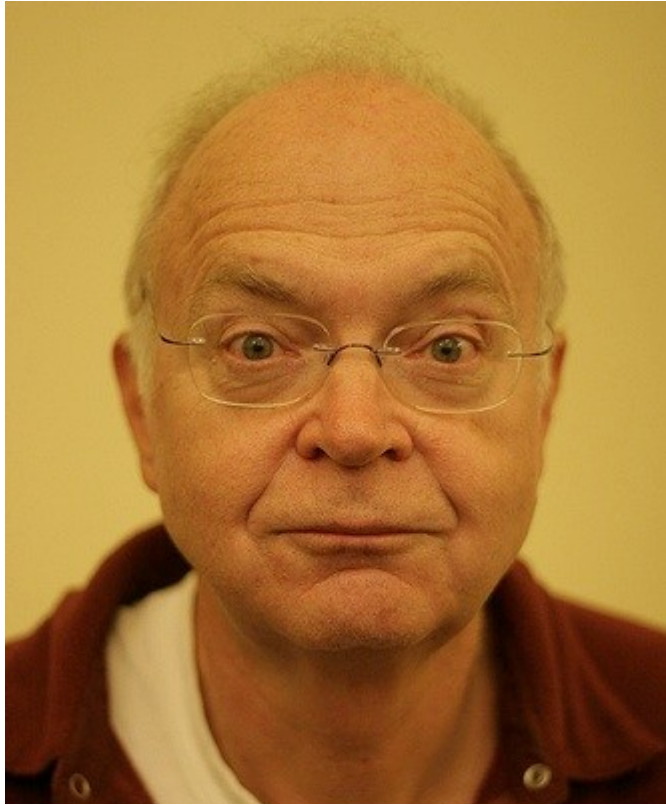
Is that all?



Landscape of Parallelization



Communicating Research



“Of equal importance to
solving a problem is the **communication**
of that **solution to others.**”
- Donald Knuth



Summary

- Parallel Computing is here to stay forever
- Make your program to run in parallel
- Innovative ideas are more valued than tiny delta ideas
- Efficient algorithms and optimizations make faster parallel programs
- GPUs are everywhere, use it wisely.
- Efficiency \propto Sustainability.



mrpraj@Gmail.com

<https://mrprajesh.co.in/pages/research.html>

Thank you.

<https://bit.ly/hpc-research-talk>

