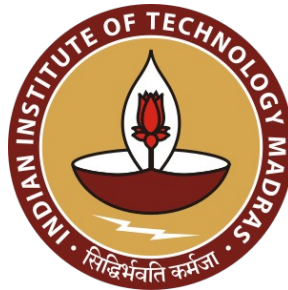


Accelerating Computation of Steiner Trees on GPUs

Rajesh Pandian M

www.cse.iitm.ac.in/~mrprajesh



Acknowledgements and disclaimer

- This is a joint work with Rupesh Nasre and N.S.Narayanaswamy
- This work evolved after the PACE Challenge 2018 on Steiner Tree [www.pacechallenge.org]
- Thanks to P100 – GPU server and PACE Lab members.
- This work is in progress / under submission.



Outline

- Steiner Tree Problem

- Example
- Properties
- Hardness

- KMB algorithm

- Challenges

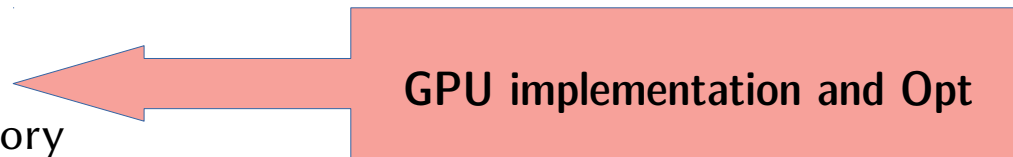
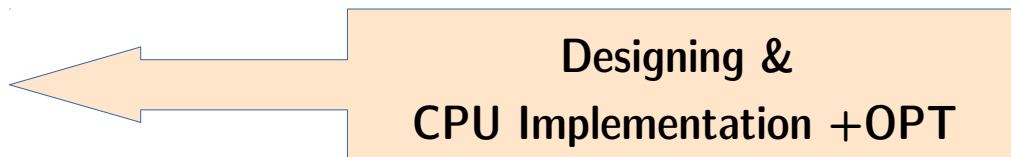
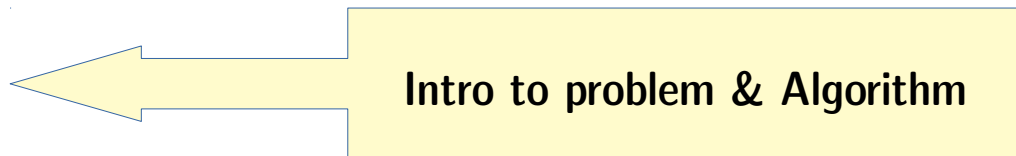
- Design Choice of KMB

- CPU Optimization

- GPU Implementation

- SSSP Optimization – Sync/Compute/memory

- Pull/Push variants



Steiner Tree Problem(STP)

Input : Undirected Graph $G(V, E, W, L)$ W is non-negative edge weights; $L \subseteq V$ terminals

Output : A tree with all terminals

Goal : Minimize the weight of the tree

- **Terminals** or terminal vertices are special vertices which must be present in the tree
- **Non-terminals** or Steiner vertices are optional vertices – generally included in tree to minimize the overall weight of the resulting tree.
- Steiner Tree - tree with all the terminals
- Applications[Hwang et. al. 92]: VLSI design, network/vehicle routing, etc.

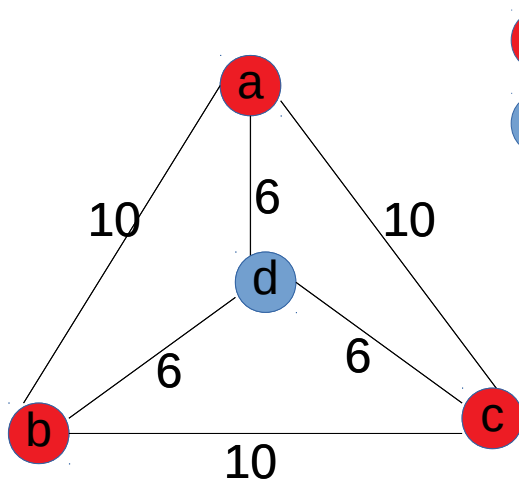


Standard Graph-theoretic notations are used $n=|V|$, $m=|E|$ and additionally $k=|L|$

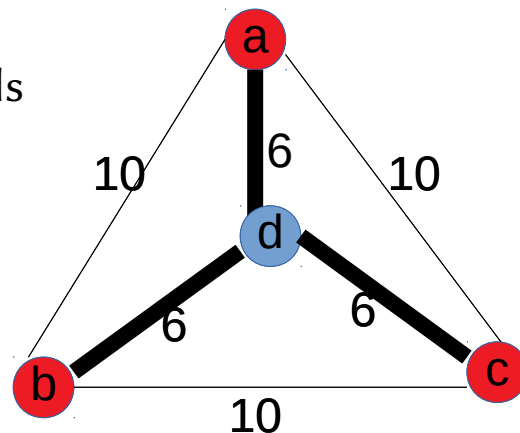
Steiner Tree Problem(STP) - Example

Input : Graph $G(V, E, W, L)$ $W:E \rightarrow \mathbb{Z}^+$ and $L \subseteq V$ terminals

Output : Connected subgraph $T'(V' \supseteq L, E' \subseteq E)$ s.t $\text{Min } W(E')$
// Minimum weighted tree with all terminals



● Terminals
● Non-terminals



Isn't MST?

Fig 1 (a)

Fig 1 (b)



Steiner Tree Problem(STP) - Example

Input : Graph $G(V, E, W, L)$ $W:E \rightarrow \mathbb{Z}^+$ and $L \subseteq V$ terminals

Output : Connected subgraph $T'(V' \supseteq L, E' \subseteq E)$ s.t $\text{Min } W(E')$
// Minimum weighted tree with all terminals

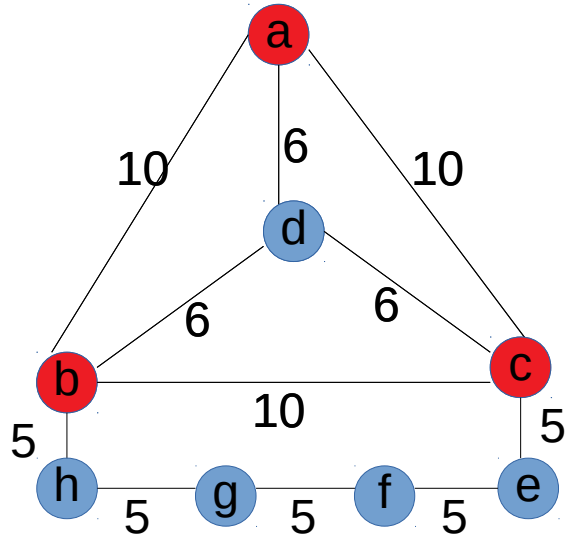


Fig 2 (a)

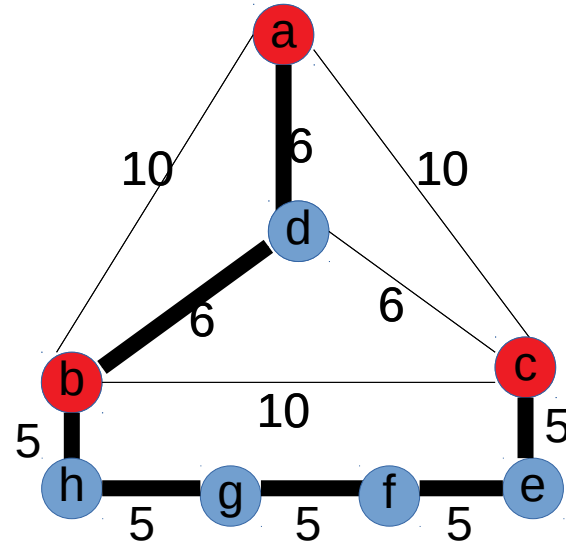


Fig 2 (b)

$|MST| = 37$
 $|OPT| = 18$



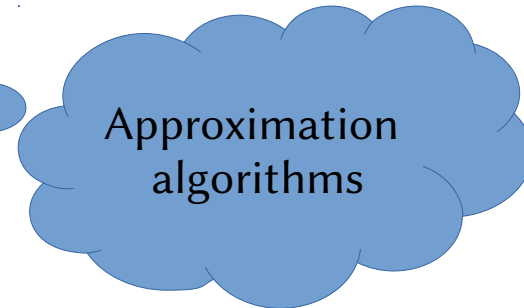
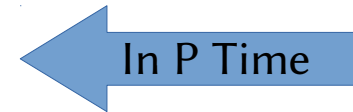
Steiner Tree Problem(STP) - Hardness

Input : Graph $G(V, E, W, L)$ $W:E \rightarrow \mathbb{Z}^+$ and $L \subseteq V$ terminals

Output : Minimum weighted tree with all terminals

Special cases

- $L = \{u, v\}$ or $k = 2$ STP=ShortestPath_In_G(u, v)
- $L = V$ or $k = n$ STP=MST(G)
- Othercases STP is NP-Hard



Standard Graph-theoretic notations are used $n=|V|$, $m=|E|$ and additionally $k=|L|$



How to deal with NP-Hardness

- No Polynomial time algorithm can find optimal solution unless $P \neq NP$
- What could be naive solutions? Enumerate all Spanning trees.
- **Approximation algorithms**
- Runs in P time
- Outputs an approximate solution with some guarantee.
 - e.g 2 or some constant, $\log n$, etc.
- There are several algorithms
 - Kou, Markowsky and Berman[KMB81]
 - Robins and Zelikovsky [RZ2000]
 - Melhorn [M88]
 - ..

$$|ALG| \leq 2 |OPT|$$



Our work

- Optimized CPU implementation for KMB Algo
- GPU optimizations for SSSP in undirected
- GPU Implementation of KMB algo.
 - Speed-up 10x (average 3x) over sequential CPU



KMB Algorithm $G(V, E, W, L)$

Phase 1

Computes the shortest distance between every pair of terminals

Phase 2

// Construct $G' = K_L$

Build a graph G' over terminals, having edge-weights corresponding to the shortest distances computed Phase 1

//Every edge in G' corresponds to a path in G

MST (G')

Phase 3

// Construct G''

For every edge in MST(G') substitute the edges with the corresponding shortest path in G

//collect all the edges & vertices of the corresponding path to construct G''

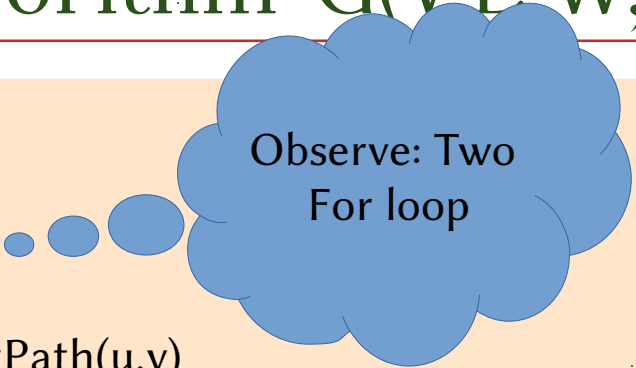
MST(G'')



KMB Algorithm $G(V, E, W, L)$

Phase 1 & 2

```
For u in L {  
  For v in L {  
    Puv = ShortestPath(u,v)  
    W'(u,v) = |Puv|  
  }  
}  
T' = MST( G'(L, -, W') )
```



Observe: Two
For loop

Phase 3

```
For (u,v) in edges of T'  
  G'' = G'' ∪ Puv  
  //Add vertices & edges of Puv  
End  
  
T' = MST ( G''( -, -, W) )
```



KMB Algorithm $G(V, E, W, L)$

Input: Graph $G(V, E, W, L)$

Output: 2-approx Steiner Tree $T(V_T, E_T)$ $V_T \supseteq L$

for $u \in L$ do

 SSSP(G, W, L, u)

 Compute W' incrementally

end

$T' = \text{MST}(G', W')$

Compute G'' and its vertices, adjList using T'

$T'' = \text{MST}(G'', W)$

return T'' ;

Single for loop but runs SSSP fully



KMB Algorithm $G(V, E, W, L)$

Input: Graph $G(V, E, W, L)$

Output: 2-approx Steiner Tree $T(V_T, E_T)$ $V_T \supseteq L$

for $u \in L$ do

parallel SSSP(G, W, L, u);

Compute W' incrementally;

end

$T' =$ **parallel** MST(G', W');

Compute G'' and its vertices, adjList;

$T'' =$ **parallel** MST(G'', W);

return T'' ;

Single for loop but runs SSSP fully

Our next work: In fact
we want to run multiple SSSP
from different source
in parallel



[Detour] SSSP : Dijkstra vs BellmanFordMoore

- Runs in time $O(n^2)$ / $O(m+n \log n)$ / $O((m+n) \log n)$
- Uses Binary/Fib Min-Heap
- At each iteration,
 - Pick up node from Q
 - RELAX'es all its neighbours
- Runs in time $O(nm)$
- No heap
- All edges are RELAX'ed at most $(n-1)$ times

For i from 1 to n-1:
For each edge (u, v) in E
RELAX(u,v, W(u,v))

In parallel setting must use Queue
In some form //it is difficult.

RELAX all edges
Launched using n threads or m



Dijkstra and its RELAX operations

INPUT: $G(V,E,W)$, src

OUTPUT: $d[]$, $p[]$

INITIALIZE-SINGLE -SOURCE (G , src)

$Q = G.V$

while(! $Q.empty()$) {

$u = \text{ExtractMin}(Q)$;

 For v in $\text{Adj}[u]$

 RELAX(u,v , W)

}

```
RELAX( $u, v, W$ ) {  
    If  $u.d + W(u,v) < v.d$  {  
         $v.d = u.d + W(u,v)$   
         $v.p = u$   
    }  
}
```

INITIALIZE-SINGLE -SOURCE(G , src)

For each v in $G.V$

$v.d = \infty$

$v.p = \text{NIL}$

}

$\text{src}.d = 0$



Source : CLRS book

Dijkstra and its RELAX operations

```
1 function Dijkstra(Graph, source):
2
3   create vertex set Q
4
5   for each vertex v in Graph:
6     distance[v] ← INFINITY
7     parent[v] ← NIL
8     add v to Q
10  distance[source] ← 0
11
```

```
12  while Q is not empty:
13    u ← vertex in Q with min distance[u]
14
15    remove u from Q
16    // only v that are still in Q
17    for each neighbor v of u:
18      newD ← distance[u] + Weight(u, v)
19      if newD < distance[v]:
20        distance[v] ← newD    //RELAX(u,v, W)
21        parent[v] ← u
22
23  return distance[], parent[]
```

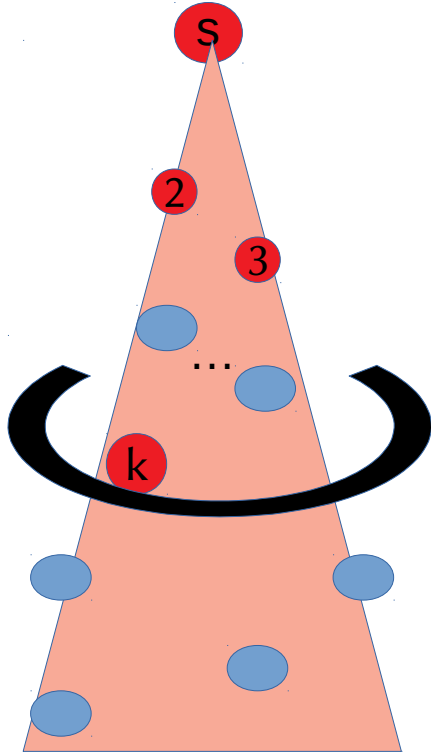


Source : Wikipedia

[coming back] CPU Implementation- Optimization

- SSSP halt optimization

Steps
of
SSSP
execution



Dijkstra Property: when a node u is picked from Q for processing then the $\text{distance}[u]$ is saturated using all the visited nodes

Halt the SSSP when all terminals are visited



GPU Implementation

- Style of implementation

Topology-driven [TD]

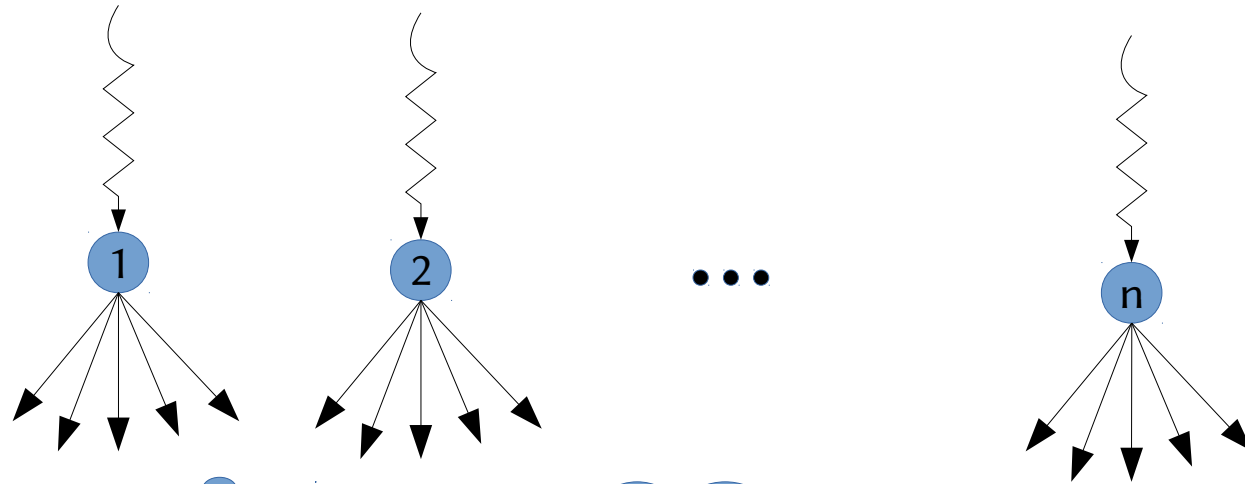
- All nodes are assumed to be active
- Operator is applied on all the nodes
- Some nodes might have no work
- Easy to implement.
- Works well on GPU

Data-driven [DD]

- Only some nodes are active
- Operates on some nodes that might have work to do
- Requires addition effort to maintain worklist(WL)
- Works well for CPU parallelism



GPU Implementation - SSSP



- n-threads
- One thread for each node
- Perform RELAX in parallel
- RELAX's its Nbrs
- Till there is no change

In some sense,
Every node/thread
pushes the data
to Nbrs



KMB Algorithm $G(V,E,W,L)$

MAIN-GPU

```
For s in L {  
  ThdsPerBlk = 512; // or 1024  
  Blks =  $\lceil n / \text{ThdsPerBlk} \rceil$ ;  
  Do {  
    INIT-KER<Blks,ThdsPerBlk>(s,  $d_s$ ,  $p_s$ , n);  
    SSSP-KER<Blks,ThdsPerBlk>(.,s,  $d_s$ ,  $p_s$ , changed, n);  
    CopyTo(DArray,  $d_s$ );           // From Device to Host.  
    CopyTo(PArray,  $p_s$ );           // From Device to Host.  
    CopyTo(hChanged, changed); // From Device to Host.  
  }while (hChanged);  
}
```

- Note we reuse $d[]$ $p[]$ across iterations
- We need the $P[]$ for know the intermediate vertices tin the shortest path



KMB Algorithm $G(V,E,W,L)$

SSSP KERNEL($..,s, d_s, p_s, done, n$)

$u = tid$ // compute tid;

if $tid < n$ {

for $v \in \text{adjacent}[u]$ { // Using CSR arrays

// Relax Operation (u, v, d_u)

$newCost = d_s[u] + W(u, v);$

$old = d_s[v];$

if $newCost < old$ then

Atomic-MIN($d_s[v], newCost$);

// Updates Parent array

if Atomic-MIN is success {

$p_s[v] = u;$

$changed = true;$

}

}}

Note

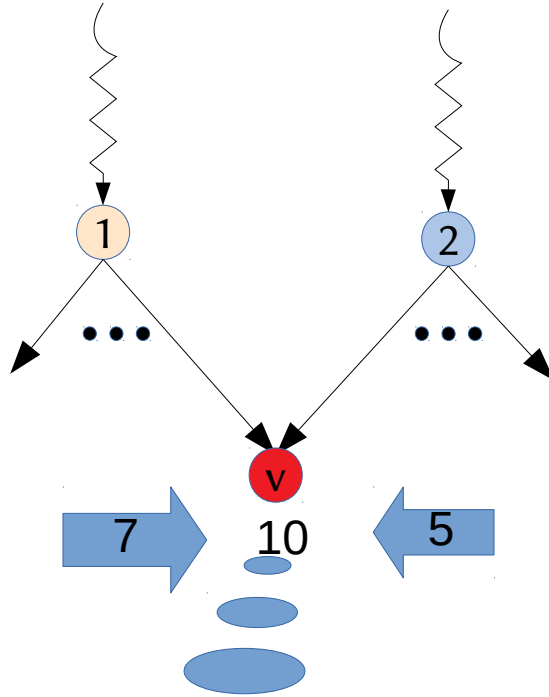
- Parent should be updated only if atomic min Success



Is it enough?



Parent update - Challenge



Suppose: 2 thread wants to update distance of its common Nbr v

<snip>

```
newCost = d[u] + W (tid, v) ;
```

```
old = d[v];
```

```
if newCost < old
```

```
    oldA=Atomic-MIN(d[v], newCost);
```

```
// AtomicMin is Success
```

```
if oldA != old {
```

```
    // Updates Parent array
```

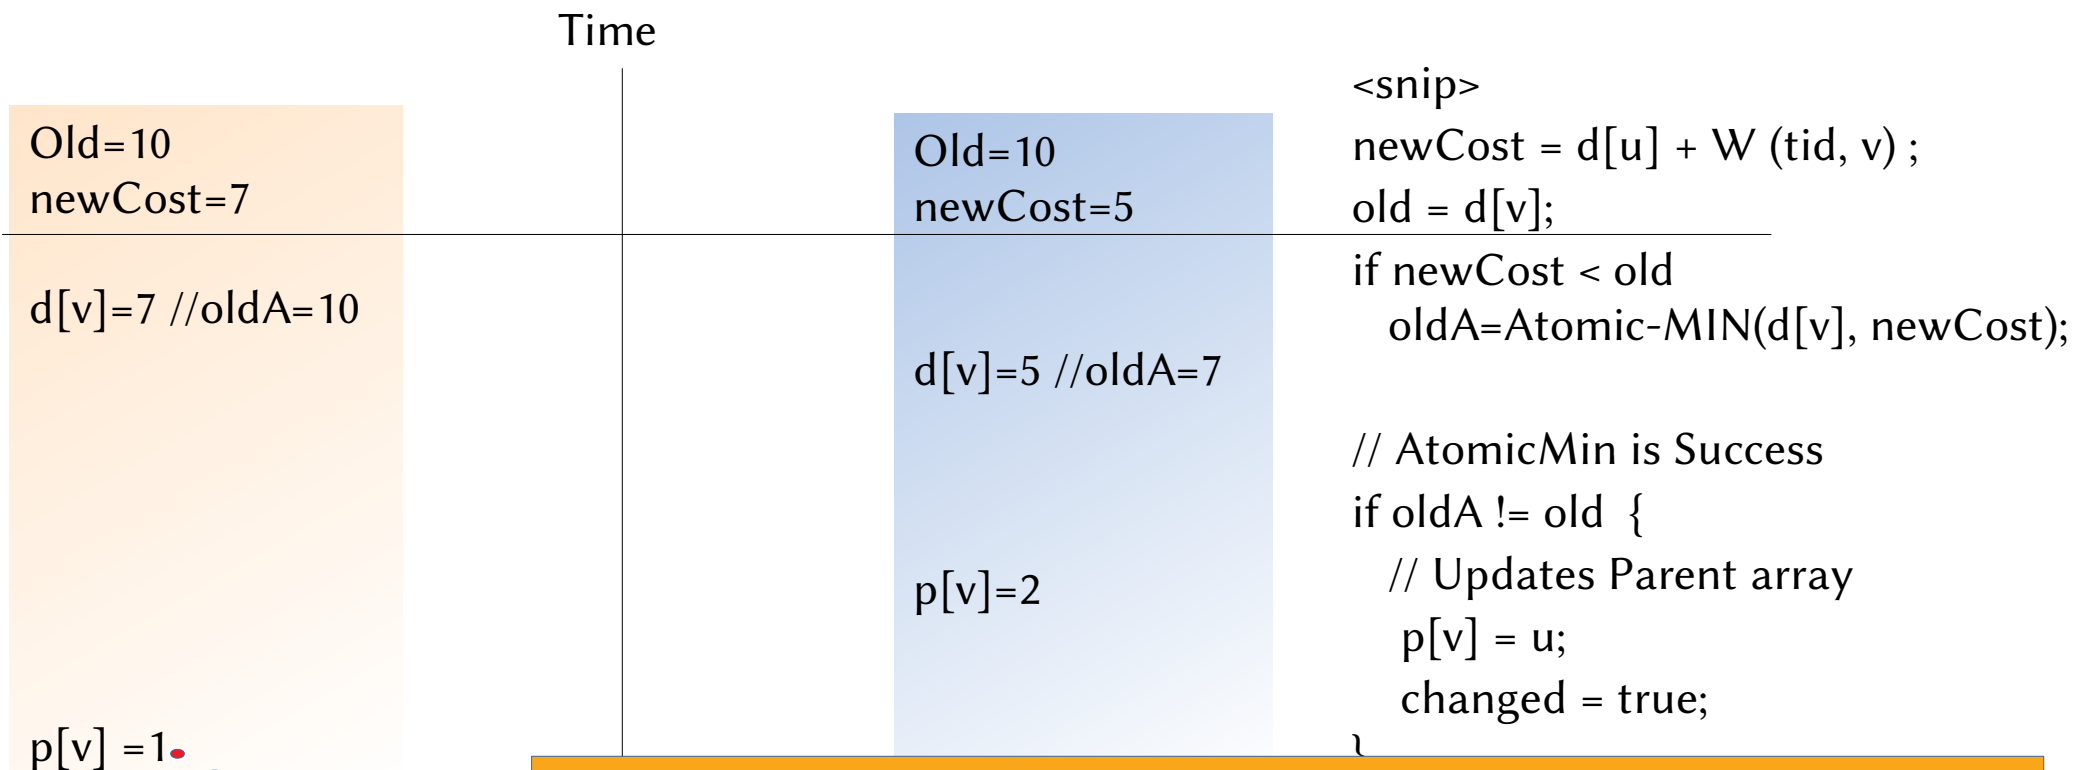
```
    p[v] = u;
```

```
    changed = true;
```

```
}
```



Parent update - Challenge



Wrong!

It is a challenge to find-out which “thread” updated distance to minimum
How update both distance and parent at the same time? Locks?



GPU Optimizations

- Synchronization
 - Push
 - Pull
- Computation
 - Data-driven
 - Edge-based
 - Controlled Computation unrolling
 - Δ
 - 2Δ
 - $t\Delta$
- Memory
 - Shared memory



Δ – max degree of the graph

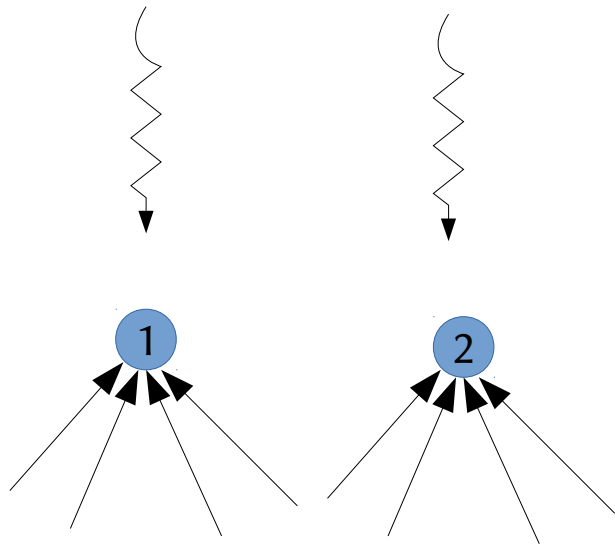
GPU Optimizations

- Synchronization
 - Push
 - **Pull**
- Computation
 - Data-driven
 - Edge-based
 - **Controlled Computation unrolling**
 - Δ
 - 2Δ
 - **$t\Delta$**
- Memory
 - **Shared memory**

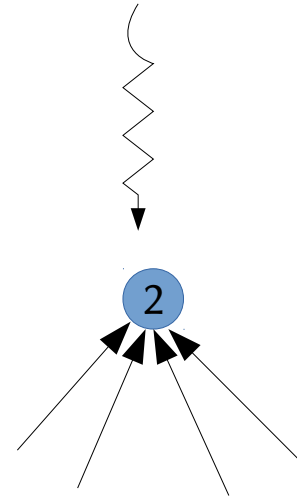


Δ – max degree of the graph

Synchronization optimization • Pull



...



```
<snip>
newCost = d[v] + W (tid, v) ;
old = d[u];
if newCost < old {
  d[u] = newCost
  p[u] = tid;
  changed = true;
}
```

No Atomics
Parent update is easy

Because, one thread is wrting on one index



Compute optimization

- Data-driven
 - Needs Worklist
 - Active/Change nodes are inserted into WL
 - Only size of WL many threads launched
 - Need synchronization while inserting nodes in WL
- Edge based optimization
 - m-threads are launched
 - RELAXes one or group of edges
 - Representation needs to be modified.
- Computation Unrolling
 - Instead of one thread doing Δ work, perform more work per thread
 - Update also neighbours of neighbours (Δ^2)
 - **Repeat the work**; Say 2 times or t times (2Δ or $t\Delta$); e.g we do pull 3 times in the kernel • 3-pull
 - When $t=3$ is performance peaked for our testcase!



Memory optimization

- Programmable shared memory can be useful
- When there are multiple reads to DRAM
- We can move data to shared memory
- For e.g. In 3-pull, we moved CSR Adj list to shared
- As the neighbours Adjlist is accessed 3 times
- Of the total 48K per block
- when using 512 threadPerBlock we have 24 words to store per thread
- Hence, if $\text{degree}(\text{node}) < 25$ we use shared, we move $\text{CSR AdjList}[\text{node}]$ to Shared
- With shared memory we achieve 24% of improvement in 3-pull



Key take-aways

- Solving Steiner Tree Problem is hard
- KMB Algorithm, an 2-approximation algorithm
- CPU Implementation has SSSP halt OPT
- SSSP with parent array parent update was challenging!
- Pull based SSSP is great for GPU KMB even without SSSP halt !



Experimental setup

t#	Graph	n	m	k	AvgDeg	MaxDeg	AvgWt	MaxWt
t1	t3-instance137.gr	97,928	1,28,632	902	2.63	14	2,486.70	2,71,369
t2	t3-instance163.gr	1,17,756	1,65,153	1,879	2.81	16	1,937.63	2,71,369
t3	t3-instance181.gr	1,35,543	2,01,803	3,033	2.98	16	1,795.73	2,58,185
t4	t3-instance183.gr	1,20,866	1,87,312	3,224	3.10	9	2,349.20	2,71,369
t5	t3-instance185.gr	66,048	1,10,491	3,343	3.35	16	1,03,010.50	2,15,10,249
t6	t3-instance187.gr	63,158	1,07,345	3,458	3.40	9	1,38,645.37	5,38,90,551
t7	t3-instance189.gr	1,72,687	2,55,825	3,902	2.96	10	2,826.30	2,71,369
t8	t3-instance191.gr	85,085	1,38,888	3,954	3.26	9	91,278.60	3,36,66,258
t9	t3-instance193.gr	17,127	27,352	4,461	3.19	4	19.71	126
t10	t3-instance195.gr	89,596	1,48,583	4,991	3.32	10	1,27,022.40	2,58,65,203
t11	t3-instance197.gr	2,35,686	3,66,093	6,313	3.11	14	2,375.55	2,71,369
t12	lin37.gr	38,418	71,657	172	3.73	4	56.17	198
t13	alue7080.gr	34,479	55,494	2,344	3.22	4	8.80	13
t14	Deezer-HR.gr*	54,573	4,98,202	3,000	18.26	420	1.00	1

Table 1. Benchmark graphs and their characteristics. Note *: t14 terminals chosen randomly with unit edge-weights



Experiments – Solutions and execution times

t#	OPT Value O	CPU Value C	GPU Value G	%deviation G vs C	%deviation G vs O	CPU Time T_C (in ms)	GPU Time T_G (in ms)
t1	11,300,427	13,038,882	13,046,669	+0.0597	+15.384	27,251.6	20,536.49
t2	13,391,485	15,075,540	15,082,504	+0.0462	+12.576	85,365.8	58,824.49
t3	20,086,478	23,054,415	23,057,792	+0.0146	+14.776	190,095	112,630.45
t4	24,998,365	28,441,249	28,446,853	+0.0197	+13.772	198,557	104,092.88
t5	793,246,106	804,364,415	804,363,424	-0.0001	+1.402	121,017	61,420.78
t6	863,275,877	869,507,927	869,511,663	+0.0004	+0.722	130,419	62,972.79
t7	40,927,523	47,966,818	47,986,233	+0.0405	+17.199	357,768	197,407.5
t8	977,020,806	989,323,251	989,322,713	-0.0001	+1.259	207,400	94,300.58
t9	184,908	197,752	198,199	+0.2260	+6.946	79,476.1	14,710.29
t10	1,406,041,806	1,424,191,280	1,424,195,265	+0.0003	+1.291	284,033	124,497.23
t11	51,655,792	58,240,533	58,251,361	+0.0186	+12.747	1,088,530	461,788.91
t12	99,560	106,937	107,297	+0.3366	+7.410	1,909.8	714.50
t13	62,449	65,529	66,290	+1.1613	+4.932	34,179.3	9,922.34
t14	-	4,629	4,595	-0.7345	-	82,829.8	8,311.58

Table 2. Comparison of solution quality and execution times (optimal value for t14 is not known; GPU time is KMBGPU's time, a pull-based implementation without any GPU optimizations.)



Experiments - Speed-up

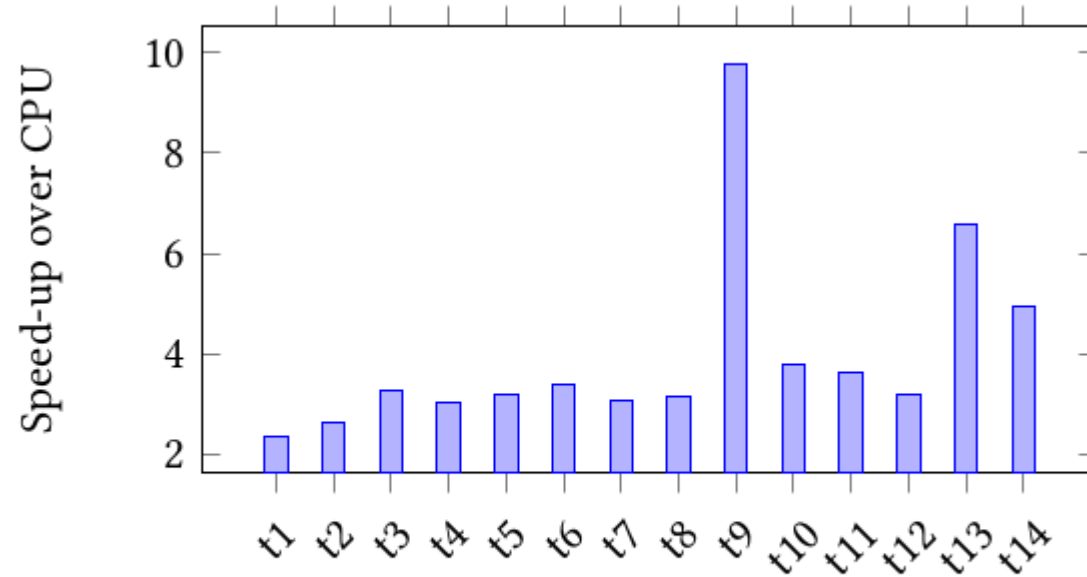


Figure 3. KMBGPU-OPT vs sequential CPU implementation



Experiments – SSSP comparision

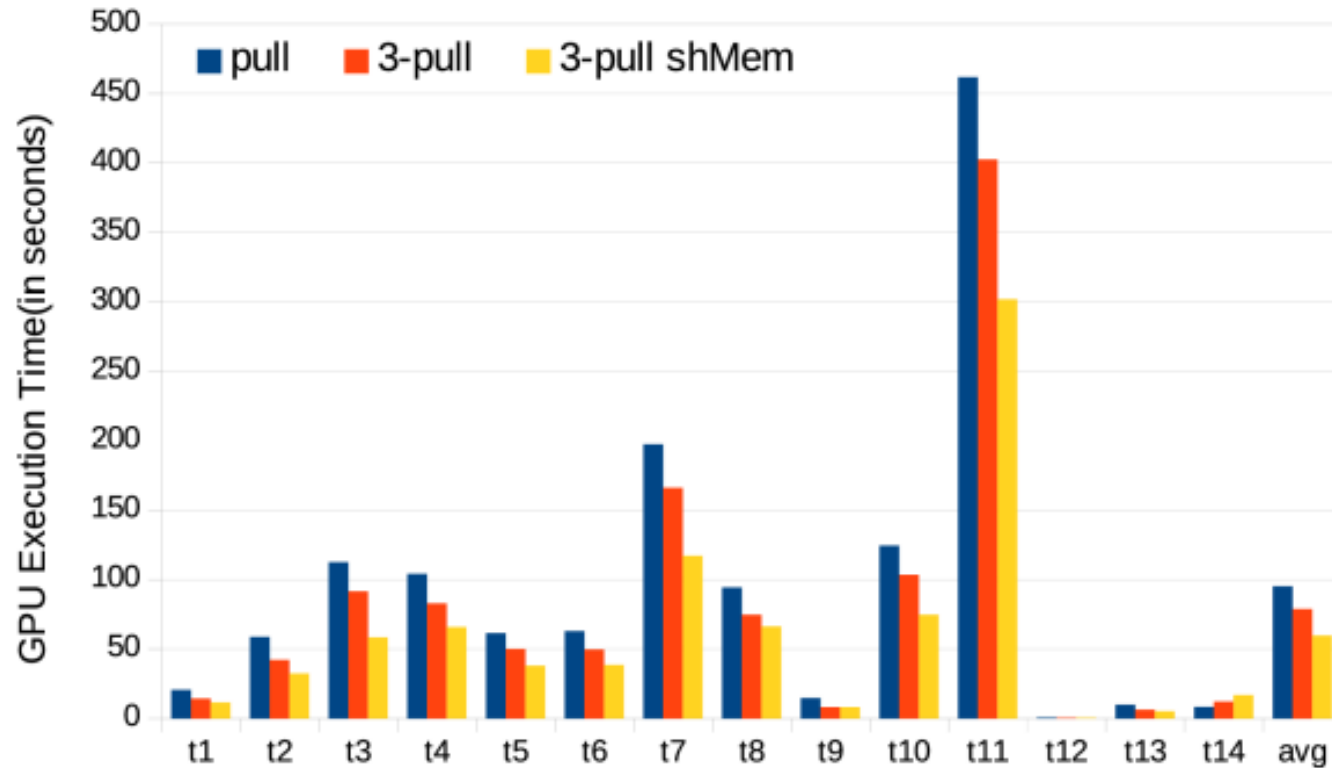
- Single SSSP
 - Using TD CSR Based
 - Data-driven Based
 - Edge-based

t#	CSR Based C	DD Based D	Edge Based E
t1	27.522	798.555	1,518.323
t2	48.634	1,472.512	2,258.170
t3	63.399	1,994.612	3,271.143
t4	31.039	966.854	2,863.490
t5	23.418	607.420	925.442
t6	15.329	262.385	744.587
t7	57.979	1,827.458	5,767.550
t8	24.292	555.236	1,605.450
t9	4.157	25.633	91.128
t10	29.921	704.242	1,461.143
t11	106.415	4,293.516	11,044.733
t12	7.518	24.018	681.824
t13	5.822	25.646	179.849
t14	4.307	17.395	5,577.330

Table 3. Comparison of CSR, Data-driven and Edge based to compute SSSP distances with source a first vertex. Note: times are in milliseconds



Comparison of GPU time with Shared memory



Questions?

Thank you,
- for your time
- for your patience

References

- Pawan Harish and P. J. Narayanan. 2007. Accelerating Large Graph Algorithms on the GPU Using CUDA. In *High Performance Computing - HiPC 2007, 14th International Conference, Goa, India, December 18-21, 2007, Proceedings*. 197–208.
https://doi.org/10.1007/978-3-540-77220-0_21
- Rupesh Nasre, Martin Burtscher, and Keshav Pingali. 2013. Data-Driven Versus Topology-driven Irregular Computations on GPUs. In *27th IEEE International Symposium on Parallel and Distributed Processing, IPDPS 2013, Cambridge, MA, USA, May 20-24, 2013*. 463–474.
<https://doi.org/10.1109/IPDPS.2013.28>
- L. Kou, G. Markowsky, and L. Berman. 1981. A fast algorithm for Steiner trees. *Acta Informatica* 15, 2 (01 Jun 1981), 141–145.
<https://doi.org/10.1007/BF00288961>
- And many...!

